



Statistics and Experimental Design using R

Sessions (Jan Kreft)

- 1.-2. Fundamental statistical concepts and R primer
3. Experimental Design
4. Probability Distributions
5. Basic single and two sample stats
6. Further single and two and more sample tests
7. Error bars, Correlation, Regression
8. Question and Answer Session
9. Assessment

Session 1-2: Fundamental statistical concepts and R primer

3

Resources


- Books
 - Michael J Crawley (2005). Statistics: an introduction using R. Wiley: Chichester, England [10 copies in Library]
 - Helmut Fritz van Emden (2008). Statistics for terrified biologists. Blackwell Publishing: Oxford [11 copies in Library]
 - Roland Ennos (2007). Statistical and data handling skills in biology. 2nd ed. Pearson Prentice Hall: Harlow, England [7 copies in Library]
 - Ruxton & Colegrave (2003). Experimental design for the Life Sciences, Oxford [20 copies in Library]
 - Fowler, Cohen & Jarvis (1998). Practical Statistics for Field Biology. 2nd. ed. Wiley [17 copies in Library]
- WebCT
 - www.weblearn.bham.ac.uk
 - Slides, data files, R scripts, Course info

2

Fundamental statistical concepts

- Why statistics?
- Significance
- p-values
- Null-hypothesis
- Effect size
- Biological importance
- The power of random processes

4



Howard C Berg (1993) Random Walks in Biology. Princeton University Press

5

Why we need statistics

- Statistics is about distinguishing real from chance effects
- A typical example is the comparison of the means of two samples
 - Compare the effect of two different treatments A and B on the body temperature of patients with fever
 - Does drug A have an effect, i.e. reduce fever?
 - 'Treatment' B could be the control placebo treatment


7

Random or not?



6


Why we need statistics



- It looks like drug A has an effect
- But randomly picking patients can also result in different samples even though there are no real differences
- Statistical analysis allows us to answer this question

8

Why we need statistics



- More precisely, statistical analysis gives a quantitative answer, it tells us how likely it is that we obtain this difference d between the means of samples A and B if they were from the same population
- In other words, how likely it is to obtain this difference d between the means of two samples if the null-hypothesis that the two treatments are equally effective were true

9

Definitions: statistical significance

- Statistical significance: An effect is called significant if it is unlikely to have arisen by chance
 - More precisely, a result is statistically significant if it is unlikely to have occurred by chance if the null hypothesis were true
 - By unlikely we usually mean a probability of less than 5%

11

Exercise

- To illustrate this point, let's create a population with known average temperature
 - Assume temperature is normally distributed
 - Mean = 37
 - SD = 1
 - `pop <- rnorm(100, mean=37, sd=1)`
 - `mean(pop)`
- Now take some samples from this artificial population of human patients
 - `sampleA <- sample(pop, 5, replace=FALSE)`
 - `mean(sampleA)`

10

Definitions: null-hypothesis

- Null-hypothesis: no effect, nothing is happening, the differences or patterns we observe are simply due to random fluctuations. For example:
 - If we compare two samples, the null hypothesis could be that the means of the two samples are the same
 - or that the variances are the same
 - or more generally that the two samples have the same distribution
 - If we compare two treatments, the null hypothesis is that there is no difference between the two treatments
 - Like any good hypothesis, it must be possible to falsify it

12

Definitions: p value

- A p value is an **estimate** of the probability that a given result (**or any more extreme result**) could have occurred by chance **if the null hypothesis were true**
- If the p value is sufficiently low, e.g. $p < 0.05$, then the
 - result is unlikely to have occurred by chance
 - result is significant
 - we reject the null hypothesis because it is unlikely to be true
- Statistics gives us a quantitative answer...
- ...but no guarantees, only an estimated probability that the null hypothesis is correct, the **p -value**

13

Exercise

- Explain the meaning of statistical significance to your neighbour
- Explain what it means when someone says “Drug A had a significant effect ($p = 0.022$)” and what it does not mean


15

Statistical versus biological ‘significance’

- A **statistically significant** difference can be **insignificant**
 - because a significant difference may be very small
 - this is about the **effect size**
 - because a significant difference may be of no (fitness) consequence
 - For example, the mean difference in height between men and women is 13 cm (UK population, age group 16-24)
 - This difference is statistically significant (as you can work out if you know sample size and variability)
 - And the effect is large (13 cm), but still...
 - ...is this difference biologically important?

14

How random walks lead to directed change




- The random walk of many particles on the microscopic scale leads to diffusion on the macroscale
 - Diffusion is directed, flux is down the concentration gradient

16

Appreciating the power of random processes

- Noise in gene regulation, say a repressor falling off by chance, can enable a cell to switch from downregulated to upregulated
 - Such random events will only happen in some of the cells, creating heterogeneity in the population
 - Need positive feedback (or double negative feedback) for such bistability



17

R primer

19

Exercise

- Think of biological processes or phenomena that are driven by chance and write your random thoughts down here now

18



What is R?

- R is a powerful software package for statistical analysis
- Originally written for teaching statistics
 - Basically an implementation of S, the standard statistical programming language (used in R and S-PLUS)
 - R comes before S
 - R was initially written by Ross Ihaka and Robert Gentleman at the Department of Statistics of the University of Auckland
- Has become a mainstream research tool
 - Users have contributed hundreds of add-on packages
 - It is very unlikely that you will encounter a statistical problem others haven't already solved in R
 - Large and active user community, so there are many to help you

20



What is R?

- And it's open-source, available for free for Win, Mac, Linux, etc.
- R is also a high-level programming language, particularly useful for numerical, technical, mathematical computing as it comes with a lot of matrix algebra functions
- Home page <http://www.r-project.org/index.html>
 - Download the latest release
 - Source code or (easier) a precompiled binary of the base system with contributed packages for a particular OS
 - E.g. for Windows, download R-2.11.1-win32.exe from a suitable mirror
 - Choose full installation to get help files in various formats (needs 65 MB disk space)
 - Otherwise you can accept the default settings

21

Some general advice

- Problem solving
 - You will run into problems when you are programming
 - Solve them by
 - reading the error message
 - reading the help file
 - make a list of what could possibly be wrong
 - check these possibilities one by one
 - break down a complex step into smaller, simpler steps
 - Divide and conquer!

23

Some general advice

- A computer is stupid: it will do exactly what you tell it to do and you need to tell it every little step, in the right order
 - Imagine writing instructions that someone can execute without having to think
 - Like describing an experimental protocol in the methods section of a paper
 - Be perfect! (Or rather learn from your errors)

22

Running R

- Once you started R (the R GUI) you will find the
 - Console (this is the command window)
 - Graphics device
 - will appear if you have created a plot
 - Script editor window
 - will appear when you create a new script or open an old one
 - here you can write your program line by line (a number of command lines to be executed in sequence)
 - it's just a text file, extension .R
 - Some terms and their meanings
 - Workspace: set of currently existing variables (objects)
 - Don't use "Save workspace..." to save your R session (commands and output), use "Save to file..."
 - History: list of previously executed commands, great for the lazy

24

R documentation and help

- `> help.start()`
 - Most useful if you don't know which command to use
 - Starts a Search Engine in your favourite browser
 - Can get this also from the R GUI Help Menu
- `> ?plot`
 - If you do know which command to use but need to check details
 - Brings up the help page on plot
- From the R GUI Help menu, you can get Manuals (in PDF)
 - An Introduction to R
 - R Reference manual

25

R basics: command line

Getting used to the command line (in red):

Prompt [Whitespace] Command [Whitespace] Return

`>` `1+2`

`[1] 3`

Results are output underneath (in blue)

The prompt tells you that R is waiting for you to enter a command, after you have entered the command, press return or enter to send it off

27

R documentation and help

- Further reading
 - Michael J Crawley (2005) Statistics: an introduction using R. Wiley: Chichester, England
 - Particularly chapter 2 and appendix and web page
 - <http://www.imperial.ac.uk/bio/research/crawley/statistics/>
 - Example data frames and further introduction to R
 - Other books
 - Various material on the R web site
 - FAQ
 - R-help mailing list

26

R basics: assignment

Assigning these results (output from command) to a variable

Variable name Assignment operator (gets) value

`> x <- 1`

`> X`

Error: object "X" not found

R is case sensitive so x and X are different variables

Variable names should start with a letter followed by letters and numbers

You can use "." or "_" instead of whitespace – but do not use whitespace since this is used to detect the end of a token ('word')

salt_concn instead of salt concn

28

R as a calculator

Arithmetic operators

- `+` add
- `-` subtract
- `*` multiply
- `/` divide
- `^` raise to the power

29

R as a calculator

Special numbers that are NOT numbers

Missing values are `NA`, Not Available

```
> x <- c(0,2,4,NA)
> x
[1] 0 2 4 NA
> mean(x)
[1] NA
> mean(x, na.rm=T) # na remove true
[1] 2
```

31

R as a calculator

How to enter numbers

```
> 10^3 # this computes 10 * 10 * 10
[1] 1000
> 1e3 # this just enters "1000" in scientific notation
[1] 1000
> 10e2
[1] 1000
> -5e3
[1] -5000
> -5e-2
[1] -0.05
> -5e-10
[1] -5e-10
```

The “#” starts a comment, any text after “#” is ignored by R

30

R as a calculator

Something totally different is `NaN`, Not any Number

Results from errors

```
> 0/0 # division by zero
[1] NaN
Like infinity, Inf
> 1/0
[1] Inf
> exp(1e200)
[1] Inf
```

32

R as a calculator

Some powers

```
> (1/27)^(1/3)
[1] 0.3333333
```

Let's do something useful: body mass index

```
> bmi <- 75/1.75^2
> bmi
[1] 24.48980
```

$$bmi = \frac{weight}{height^2}, \text{ units } \left[\frac{kg}{m^2} \right]$$

33

Exercise: R as a calculator

- Calculate and assign result into variable with name z
 - Square root of 2
 - Sine of 2π
 - Natural logarithm of 10 and 10^{1000}
 - 2^{10}
 - e^{10}
 - Calculate the values of the following functions

$$\frac{1}{9}, \left(\frac{1}{9}\right)^2, \sqrt[3]{27}, \left(\frac{1}{9}\right)^{\frac{1}{2}}, \sqrt{\frac{1}{9}}$$

35

R as a calculator

- All the usual functions exist (log, exp, sin, cos, tan, sqrt, etc.)
- ```
> log(10)
```
- ```
[1] 2.302585
```

 #so log() is the natural logarithm
- ```
> exp(1)
```
- ```
[1] 2.718282
```

 # this is e
- ```
> pi
```
- ```
[1] 3.141593
```

 # π is known as pi
- ```
> sin(pi)
```
- ```
[1] 1.224606e-16
```

 # should be zero and is fairly close
- ```
> cos(pi)
```
- ```
[1] -1
```

34

R basics: writing functions

- Writing your own functions and understand what a function is
- Example

```
• add <- function(arg1, arg2) {
•   result <- arg1 + arg2
•   result
• }
```

- NB: Write this as a script in the script editor, then execute
- Use like this

```
• > s <- add(1,2)
• > s
• [1] 3
```

36

R basics: data structures

- Scalar
 - single value of a physical quantity (product of numerical value and physical unit)
 - e.g. 250 μM oxygen concentration in lake water
- Vector
 - series of scalars
 - e.g. oxygen concentration gradient along a transect
 - e.g. oxygen concentration in a number of lakes
- Matrix (2D Array)
 - table of scalars (e.g. a field of scalars, spatial)
 - e.g. oxygen concentration field over lake surface (2D) or including depth (3D matrix/array)

37

R basics: constructing vectors

Many ways to make vectors

Simplest is the colon “:”

```
> x <- 5:1  
[1] 5 4 3 2 1
```

```
> x <- 0:4  
[1] 0 1 2 3 4
```

Versatile is the concatenate function “c”

```
> x <- c(0,2,-1,pi,10)  
[1] 0.000000 2.000000 -1.000000 3.141593 10.000000  
> x <- c(0:5,10:15)  
[1] 0 1 2 3 4 5 10 11 12 13 14 15
```

39

R basics: data structures

- Dataframes
 - O₂ concn as response variable, lake depth, pH, season, nutrient availability as (potential) explanatory variables
 - This is how you should organize most of your results
- Lists
- Tables
- Timeseries

38

R basics: constructing vectors

Sequences with finer control than “:” (colon operator) with seq

```
> x <- seq(0,2,.25)  
[1] 0.00 0.25 0.50 0.75 1.00 1.25 1.50 1.75 2.00
```

```
> length(x)
```

```
[1] 9
```

```
> x <- seq(0,2,length=11)
```

```
[1] 0.0 0.2 0.4 0.6 0.8 1.0 1.2 1.4 1.6 1.8 2.0
```

Repeating elements with rep

```
> rep(1:2,4)
```

```
[1] 1 2 1 2 1 2 1 2
```

40

R basics: working with vectors

- Functions also work on vectors (element-wise)!
- `> x <- 1:10`
- `> x`
- `[1] 1 2 3 4 5 6 7 8 9 10`
- `> x2 <- x^2`
- `> x2`
- `[1] 1 4 9 16 25 36 49 64 81 100`
- `> sqrt(x2)`
- `[1] 1 2 3 4 5 6 7 8 9 10`

41

Exercise: vectors

- Make these vectors
 - 0.00 0.25 0.50 0.75 1.00
 - 0 7
 - 1 2 3 1 2 3 1 2 3
 - Make a vector x containing numbers 0,1,2,...,10
 - Calculate vectors y according to these equations
$$y = x^2, \quad y = \sqrt{x}, \quad y = \frac{1}{x}, \quad y = \frac{1}{x^2}$$
 - Make a vector x containing 100 elements from $-\pi$ to π (equally spaced)
 - Calculate vectors y according to these equations
$$y = 2\sin(x), \quad y = \cos(2x)$$

43

R basics: working with vectors

- `> p <- seq(-2*pi,2*pi,pi/2)`
- `> p`
- `[1] -6.283185 -4.712389 -3.141593 -1.570796`
`0.000000 1.570796 3.141593 4.712389`
`6.283185`
- `> sin(p)`
- `[1] 2.449213e-16 1.000000e+00 -1.224606e-16`
`-1.000000e+00 0.000000e+00 1.000000e+00`
`1.224606e-16 -1.000000e+00 -2.449213e-16`

42

R basics: vector subscripts (indices)

- Accessing blocks of vectors with subscripts
 - Distinguish index from value of a vector element
 - You access some elements of vectors rather than all using `[]`
 - `v <- -50:50` # make simple example vector
 - `mean(v)` # this uses all elements of v
 - `v[1]` # first element
 - `v[40:60]` # a section in the middle
 - `v[length(v)]` # last element since `length()` gives number of elements
 - `v[(length(v)-9):length(v)]` # use the last 10 elements (maybe you want to select the final stage)

44

R basics: vector subscripts (indices)

- Something special: you can select elements of a vector that fulfil a logical condition
 - `v[v > 0]` # all elements for which condition is true will be selected
 - `v[v == 50]` # exactly equal, may give no hit
 - `v[v == mean(v)]` # also unlikely to get a hit

45

R basics: matrix subscripts (indices)

- Accessing blocks of matrices with subscripts using `[]`
 - Arrays with 2 dimensions (2D array or 2D matrix)
 - 1st dimension: rows
 - 2nd dimension: columns
 - That's how we read newspapers
 - Accessing blocks of elements in an 2D array
 - `A[1,2]`
 - Read like this: [rows , columns]
 - `A[,1]`
 - Empty means everything (all rows, or all columns)

47

Exercise: vector subscripts (indices)

- Make a vector from 0 to 99
 - Select these elements
 - first
 - last
 - 11 in the middle
 - the last 50
 - Select elements based on logical conditions
 - all elements smaller than 10
 - all elements smaller and equal to 10
 - all elements equal to 10
 - all elements equal to 100
 - all elements equal to 0.1

46

R basics: matrix subscripts (indices)

- Construct our example array from a vector
 - `> v <- 0:15`
 - `> v`
 - `[1] 0 1 2 3 4 5 6 7 8 9 10 11 12`
`13 14 15`
 - `> A <- matrix(v,4,4)`
 - `> A`
 - `[,1] [,2] [,3] [,4]`
 - `[1,] 0 4 8 12`
 - `[2,] 1 5 9 13`
 - `[3,] 2 6 10 14`
 - `[4,] 3 7 11 15`

48

Exercise: matrix subscripts

- Now play with this matrix, selecting
 - the entire first row
 - the entire first column
 - the block in the middle
 - the last two entire columns
 - all values larger than 10

49

R basics: dataframes

- Dataframes: THE data structure for your results
- Create table of results in Excel (that's easier)
 - “Save as” “Text file (Tab delimited) (*.txt)”
 - Make sure you know where you saved the file, and its name
 - Let's call the file “worms.txt”
- In R, you can use the menu (File > Change dir...) to select the directory where you have saved the file, makes entering the file name easier
 - Read this txt file with `read.table()` like this
 - `worms <- read.table("worms.txt", header=TRUE, row.names=1)`
 - File name must be correct and exist where R is looking & put in quotes
 - The first row is the header, containing the names of the variables
 - The first column contains the row names (don't treat Field.Name as variable)
 - Variable names must be ONE word (no blanks)
 - Missing values should be denoted NA (no blanks)

This is a string
Always put in "quotes"

51

R basics: dataframes

You can make such dataframes in Excel!

	Explanatory variables					Response variable
Field.Name	Area	Slope	Vegetation	Soil.pH	Damp	Worm.density
Nashs.Field	3.6	11	Grassland	4.1	F	4
Silwood.Bottom	5.1	2	Arable	5.2	F	7
Nursery.Field	2.8	3	Grassland	4.3	F	2
Rush.Meadow	2.4	5	Meadow	4.9	T	5
Gunness.Thicket	3.8	0	Scrub	4.2	F	6
Oak.Mead	2.1	7	Grassland	3.9	F	2
Church.Field						
Ashurst						
The.Orchard						
Rookery.Slope						
Garden.Wood						
North.Gravel	3.3	1	Grassland	4.1	F	1
South.Gravel	3.7	2	Grassland	4	F	2
Observatory.Ridge	1.8	6	Grassland	3.8	F	0
Pond.Field	4.1	0	Meadow	5	T	6
Water.Meadow	3.9	0	Meadow	4.9	T	8
Cheapside	2.2	8	Scrub	4.7	T	4
Pound.Hill	4.4	2	Arable	4.5	F	5
Gravel.Pit	2.9	1	Grassland	3.5	F	1
Farm.Wood	0.8	10	Scrub	5.1	T	3

All values of the same variable must be in the same column!

50

R basics: dataframes

- Assigning the data to the variable “worms”, the dataframe will be known by this name
 - `worms <- read.table("worms.txt", header=TRUE)`
 - `worms`
 - `worms <- read.table("worms.txt", header=TRUE, row.names=1)`
 - `# now the first column is treated as row label, not explanatory variable`
 - `worms`
 - `class(worms) # yes, it is a dataframe`
 - `names(worms) # lists the names of the variables in the dataframe`
 - `plot(worms) # plots all variables against all others, pH seems to be the most interesting`

52

R basics: dataframes

- `plot(Soil.pH,Worm.density)` # error message, the variables are not known!
- `attach(worms)` # this makes the variables inside the dataframe visible to the outside world
- `plot(Soil.pH,Worm.density)` # now it works
- `summary(worms)`
- # now let's try subscripting to select parts of the dataframe
- `worms[,4]` # all rows, column 4
- `worms[Soil.pH >= 5,]` # those rows with pH >= 5, all columns

53

R basics: dataframes

- # these commands do the same but are less easy to read
- `ws <- worms[order(Soil.pH),]` # all in one line
- `ws <- worms[order(worms[,4]),]` # the same
- Remove variables so we can use the same names again without creating confusion
- `detach(worms)` # undo attach
- `rm(worms)` # rm for remove

55

R basics: dataframes

- Sorting and ordering (use `order()` for dataframes)
- `sort(Soil.pH)` # don't use sort on dataframes, it will only sort one column, leaving the rest of the row as it was
- `i <- order(Soil.pH)` # use order instead, order returns subscripts (indices)
- `i` # let's have a look at the indices `i`
- `Soil.pH` # unordered
- `Soil.pH[i]` # ordered
- `worms[i,]` # now the whole dataframe in order

54

Exercise: dataframes

- Example: Enzyme activity as a function of substrate concn
 - Download file KMINDOL.xls from WebCT
 - Open file in Excel
 - Inspect data: are there variable names in the header, are there row names, and if so, in which column?
 - Save as tab-delimited text file KMINDOL.txt
 - Read text file into R dataframe
 - Examine variable names and dataframe content
 - Attach dataframe
 - Plot the response variable versus the explanatory variable

56

R: plotting

- Plot a function to see its shape

```
x <- seq(-2*pi,2*pi,length=101) # make vector x
y <- sin(x) # calculate vector y as a function of x
plot(x,y)
plot(x,sin(x)) # another way of doing the same
```
- The result is OK, but for publication quality this default plot is not good enough, we want lines not points (type="l"), y-axis labels horizontal (las=1), etc.

```
plot(x,y, las=1, xlab="x", ylab="sin(x)", type="l",
     main="Trigonometric functions")
```
- Many more parameters can be set within the plot() call

57

R: plotting

- Low-level plotting functions add to an existing plot
- Let's plot sin(x) again

```
plot(x,y, las=1, xlab="x", ylab="f(x)", type="l")
```
- Now that we have a plot, we can add

```
lines(x,cos(x),col="red")
points(x, cos(x)+rnorm(length(x), mean=0, sd=0.1),
       col="blue") # cos(x) with some noise added
title("Trigonometric functions")
text(0.5*pi+pi/4, sin(0.5*pi), "sin(x)")
```
- Example low-level plotting functions: lines, points, text, title, legend, grid, arrows, etc.

59

R: plotting

- Distinguish high-level from low-level plotting functions
 - High-level plotting functions create a new plot, complete with axes, ticks, labels etc.
 - These are difficult to change afterwards!
 - Make plot behave by setting graphics parameters before the call to plot, e.g.

```
par(las=1)
```
 - Or set them within the high-level plotting function, e.g.

```
plot(x,y, las=1, ...)
```
 - Example high-level plotting functions
 - plot(), qqplot(), hist(), image(), contour(), wireframe(), cloud(), etc.

58

R: plotting

- Here is a list of options for plot() that are commonly used (there are far more). Many can also be used in other plotting functions

```
xlim=c(0,80), ylim=c(0,1) # set x and y axis limits
log="x", log="y", log="xy" # logarithmic axes
asp=1 # aspect ratio of axes (y/x), use 1 for square axes
type="whatever" # l for lines, p for points, b for both, n for nothing, o for both overplotted, s for stairs
col="yellow" # never use yellow, it is very hard to see
lwd=1 # line width
lty=1 # line type: 1 solid, 2 dashed, etc up to 6, more control with "8414"
pch=1 # plot character, can use integers in range 0:25 or (easier) a single character in quotes, e.g. pch="+", pch="*", pch="?"
```


60

Exercise: plotting

- Plot all functions into one plot
 - Make vector for x axis, from 0 to 2 (make 100 points for smooth curve)
 - Plot $y=x^2$ versus x as a red line
 - Plot $y=\sqrt{x}$ versus x as a blue line
 - Plot $y=x$ versus x as a black line

61

Vector graphics versus pixel graphics (bitmaps)



63

From plot window to publication

- eps: encapsulated postscript
- postscript is a page description language, a programming language high end printers can understand (interpret)
 - basically a text file containing commands for the printer
 - put this there
- encapsulated postscript is for single objects rather than a complete printed page/document
- a bit like pdf, also from Adobe
- used by publishers
- ideal for submission of vector graphics to journals
- (for bitmaps, use tiff)

62

From plot window to PowerPoint

- Right click on R graphics window and select
 - Save as postscript...
 - Select encapsulated postscript as file type
 - Choose a file name, extension eps
- Open file in GSview (GhostScript view) to display eps file
 - GSview can convert to pdf, print to printer, etc.
 - You need to install GhostScript and GSview
 - <http://pages.cs.wisc.edu/~ghost/>
- Zoom to the final size you want to use in PowerPoint
 - Do not change the size in PowerPoint, quality will suffer
- Copy the image window (Control-C)
- Paste into PowerPoint

64

From plot window to Word

- Save eps file as before
- Insert eps file in Word
 - Insert picture from file
 - Word cannot display eps, instead it displays a crude bitmap of the eps, a sort of thumbnail tiff image stored inside the eps file
 - You can print this file to a postscript printer no problem
 - You cannot print to normal cheap printers
 - But you can save as pdf and print that

65

List of some useful R functions

- Central tendency
 - mean()
 - median()
- Other useful functions
 - sum()
 - prod()
 - length()
 - sort()
 - order()
 - min()
 - max()
 - abs()
- Measures of variation
 - range()
 - quantile() (percentiles)
 - boxplot()
 - var()
 - sd()
 - mad() (median absolute deviation)

67

From plot window to Vector graphics software

- For further editing of the figure such as annotations with arrows, text, equations
- Save eps file as before
- Import or open the eps file in a good Vector graphics program that can interpret the eps format as vector graphics (rather than producing a bitmap from it, which will not be editable)
 - Corel Draw can import and export eps
 - Inkscape can import pdf (R can export figure as pdf)

66